

CSE 417T

Introduction to Machine Learning


Lecture 7

Instructor: Chien-Ju (CJ) Ho

Recap

VC Generalization Bound



- VC Bound: $E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{d_{vc} \frac{\ln N}{N}}\right)$
- The performance of your learning, i.e., $E_{out}(g)$, depends on
 - How well you fit your data ($E_{in}(g)$)
 - How well your $E_{in}(g)$ generalizes to $E_{out}(g)$



Bias-Variance Decomposition


$$\bullet \mathbb{E}_D[E_{out}(g^{(D)})] = \mathbb{E}_{\vec{x}} \left[\overset{\text{Bias}(\vec{x})}{(\bar{g}(\vec{x}) - f(\vec{x}))^2} \right] + \mathbb{E}_{\vec{x}} \left[\mathbb{E}_D \left[\overset{\text{Var}(\vec{x})}{(g^{(D)}(\vec{x}) - \bar{g}(\vec{x}))^2} \right] \right]$$

- The performance of your learning, i.e., $\mathbb{E}_D[E_{out}(g^{(D)})]$, depends on
 - How well you can fit your data using your hypothesis set (**bias**)
 - How close to the best fit you can get for a given dataset (**variance**)




Learning Curves

Simple Model




Complex Model




Learning Curves

VC Analysis

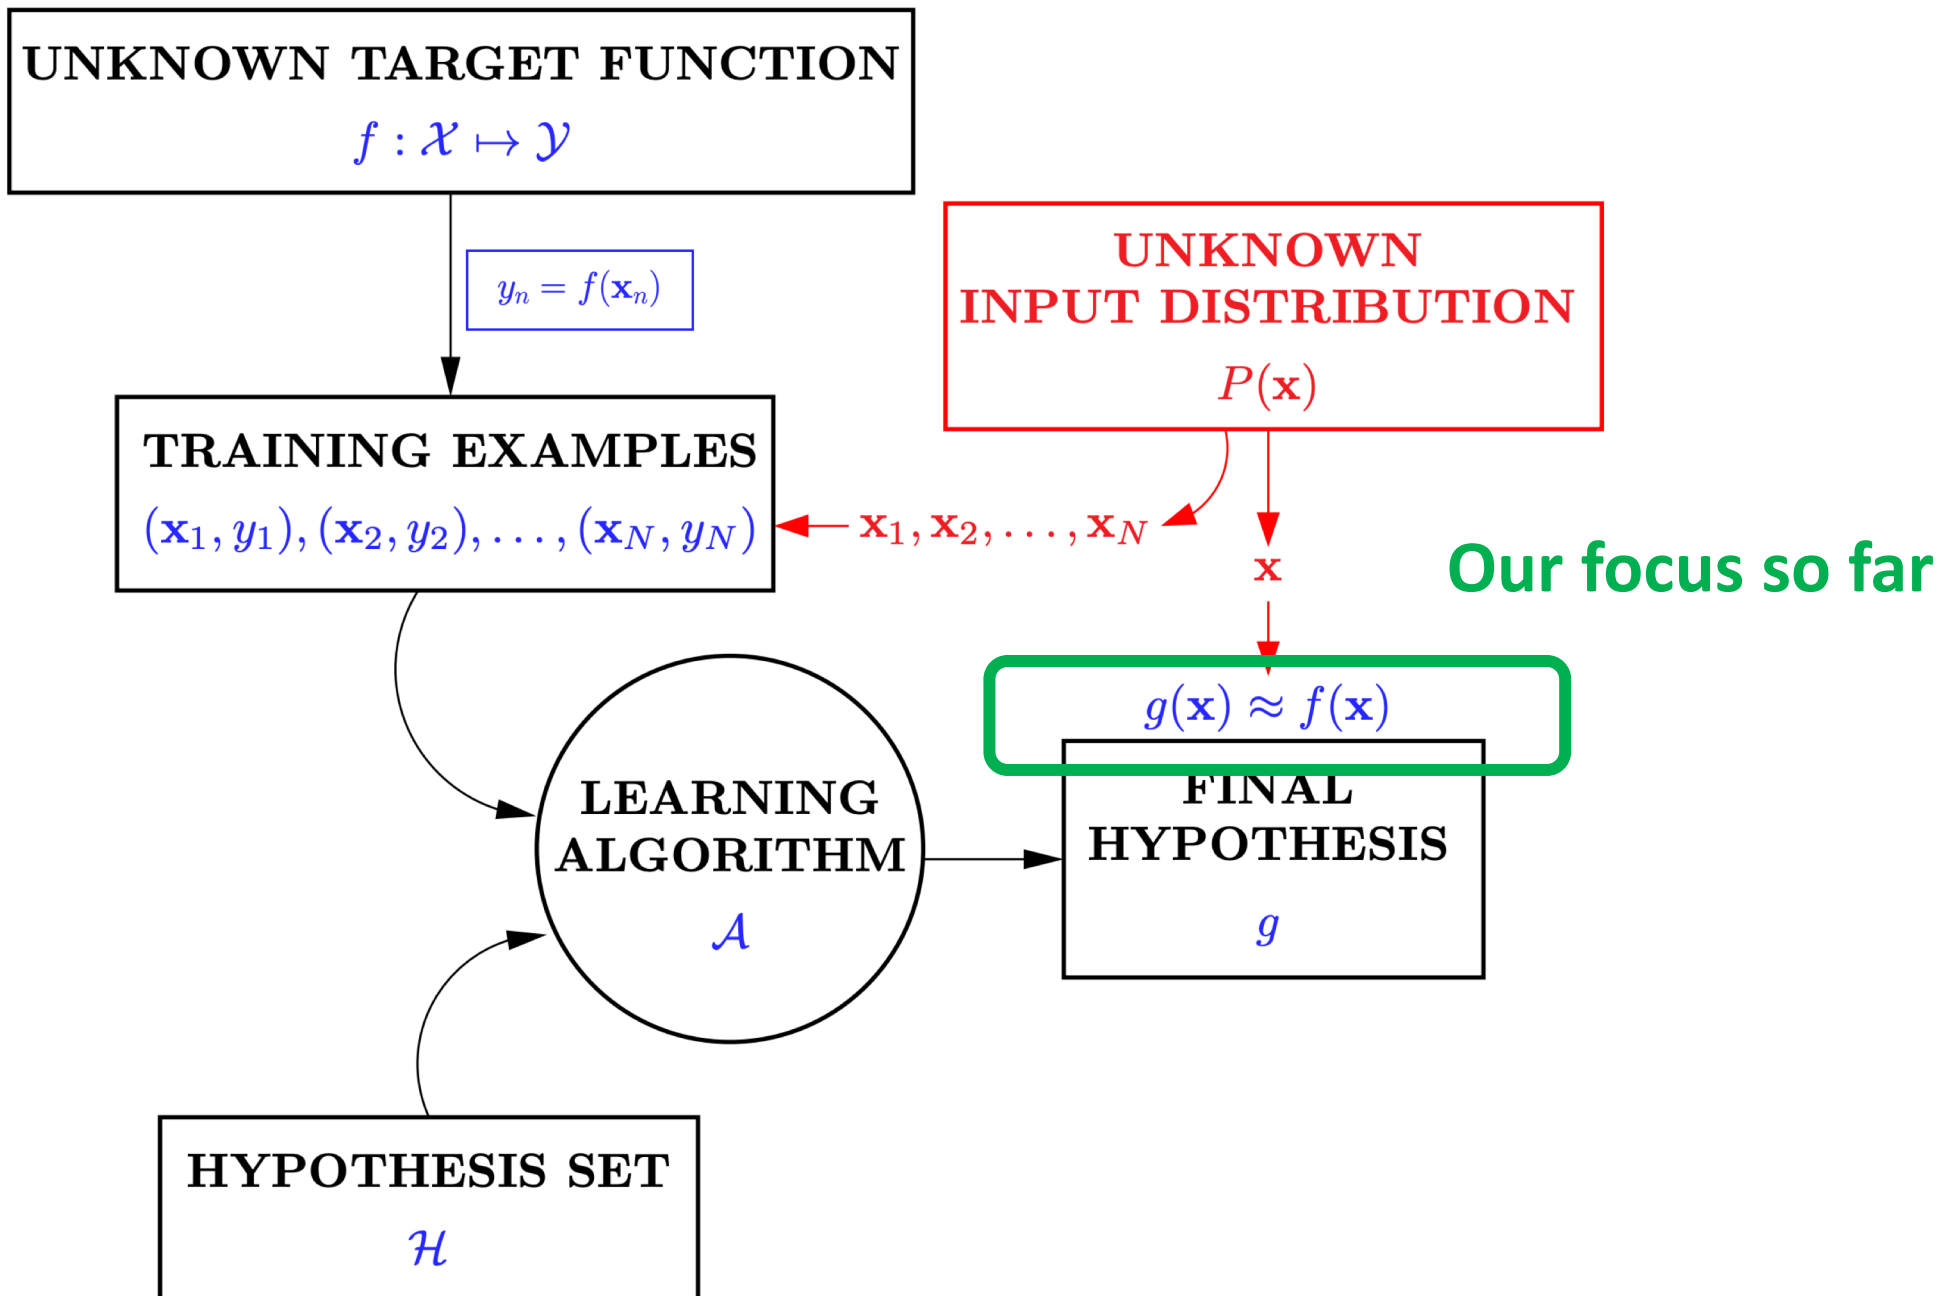


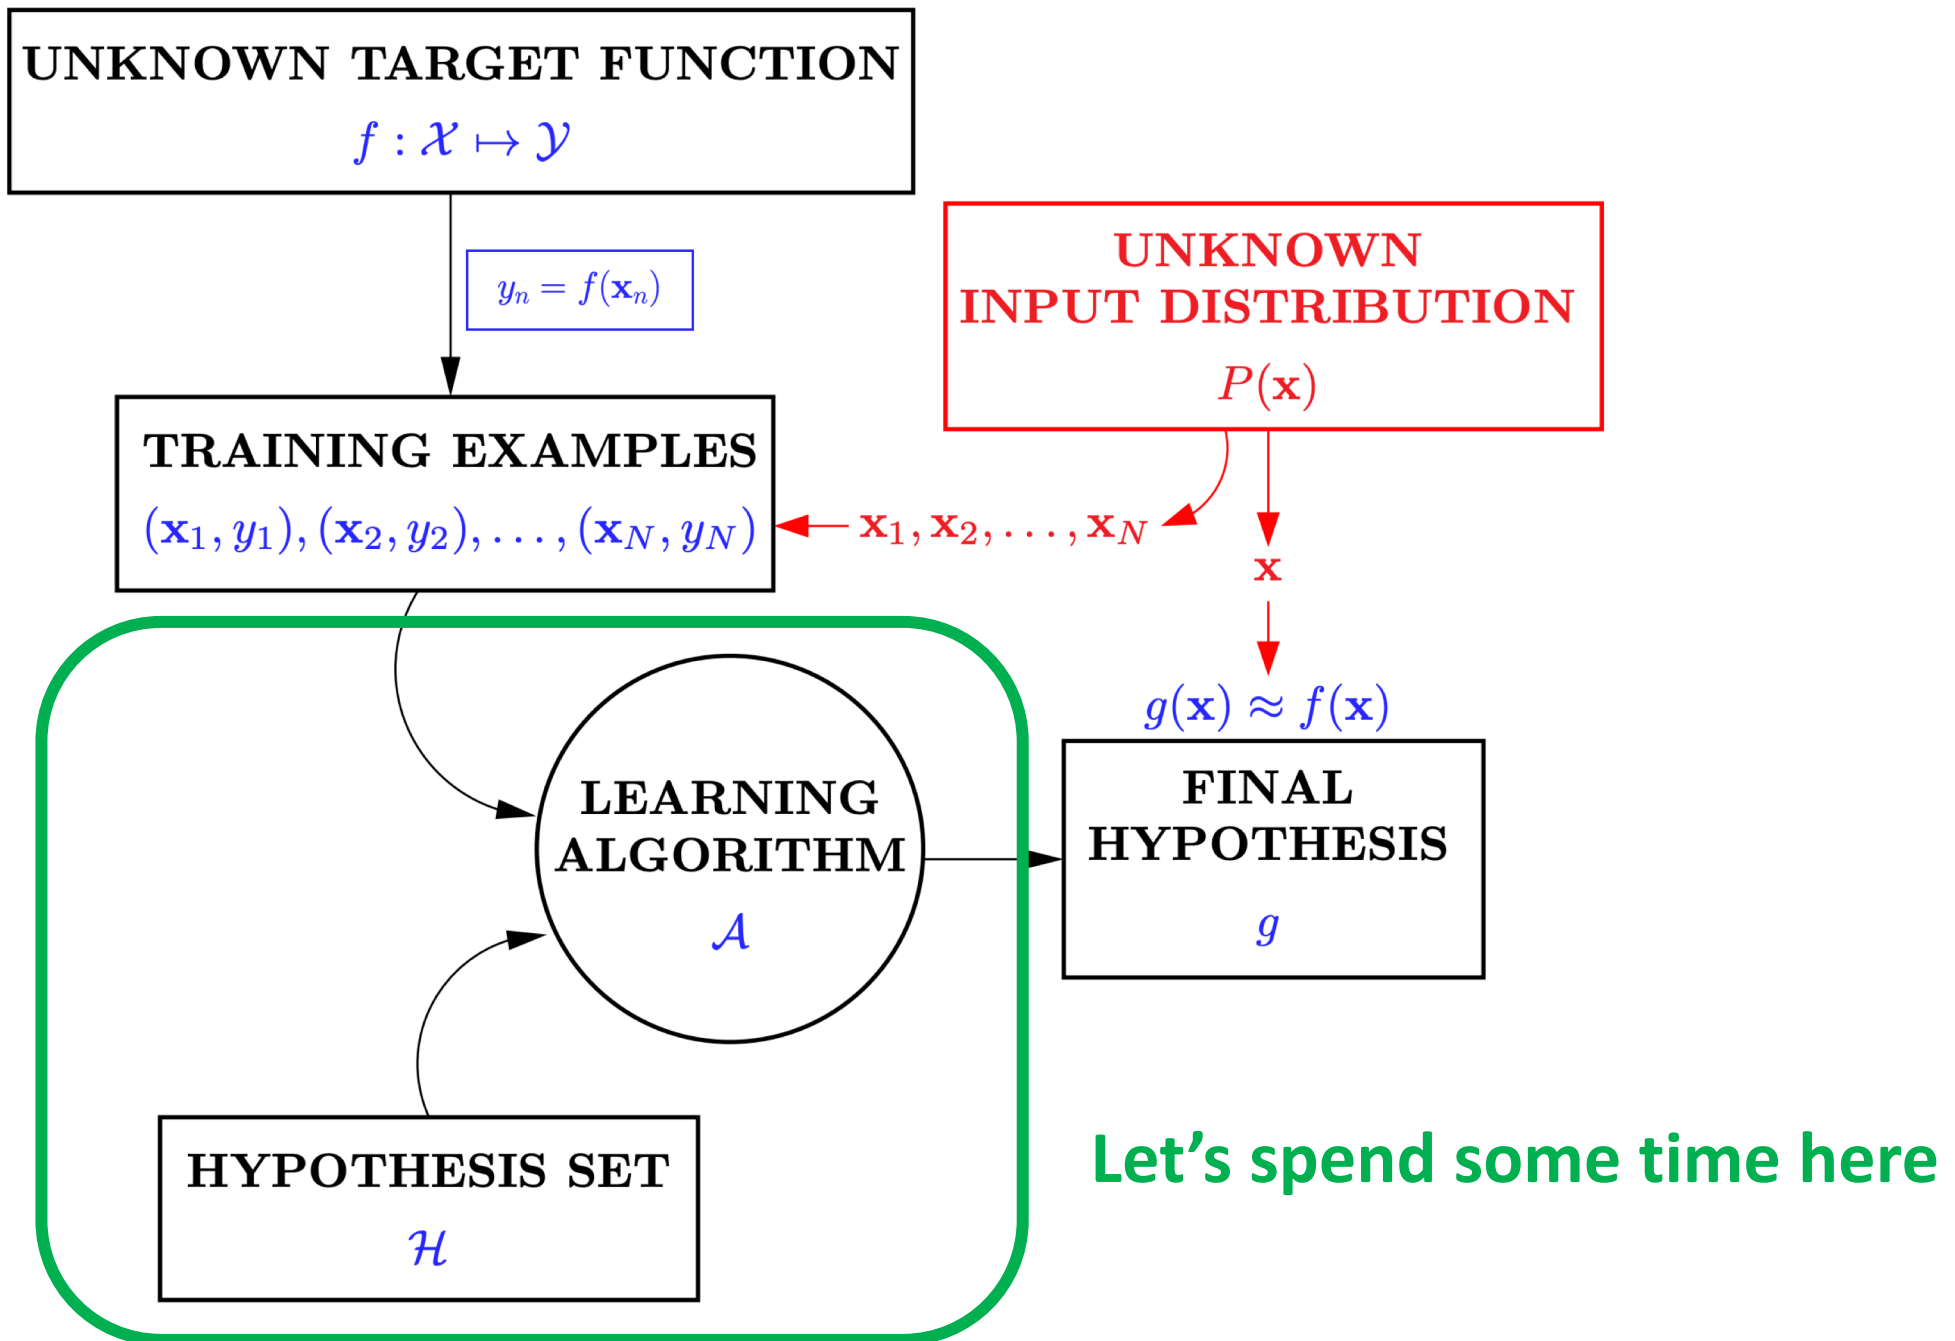
Bias-Variance Analysis



Brief Lecture Notes Today

The notes are not intended to be comprehensive. They should be accompanied by lectures and/or textbook.
Let me know if you spot errors.





Linear Models

Linear Models

This is why it's called linear models

- H contains hypothesis $h(\vec{x})$ as **some function of $\vec{w}^T \vec{x}$**

	Domain	Model
Linear Classification	$y \in \{-1, +1\}$	$H = \{h(\vec{x}) = \text{sign}(\vec{w}^T \vec{x})\}$
Linear Regression	$y \in \mathbb{R}$	$H = \{h(\vec{x}) = \vec{w}^T \vec{x}\}$
Logistic Regression	$y \in [0,1]$	$H = \{h(\vec{x}) = \theta(\vec{w}^T \vec{x})\}$

$$\theta(s) = \frac{e^s}{1 + e^s}$$

- Linear models:
 - Simple models => Good generalization error
- Reminder:
 - We will **interchangeably use h and \vec{w}** to represent a hypothesis in linear models

Algorithms?

- Goal of the algorithm:
 - Find $g \in H$ such that $g \approx f$
 - Define **error measures** to quantify $g \approx f$
 - Find $g \in H$ that minimizes $E_{out}(g)$
- Recall on the error measure
 - Often focus on point-wise error $e(h(\vec{x}), f(\vec{x}))$
 - Binary error for classification
 - Squared error for regression
 - $E_{in}(h) = \frac{1}{N} \sum_{n=1}^N e(h(\vec{x}_n), f(\vec{x}_n))$
 - $E_{out}(h) = \mathbb{E}_{\vec{x}}[e(h(\vec{x}), f(\vec{x}))]$

Algorithms?

- Goal of the algorithm: Find $g \in H$ that minimizes $E_{out}(g)$
- VC Bound: $E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{d_{vc} \frac{\ln N}{N}}\right)$
- Common algorithms:
 - $g = \operatorname{argmin}_{h \in H} E_{in}(h)$
 - Works well when the model is simple (generalization error is small)
 - Will focus on this in the discussion of linear models
 - $g = \operatorname{argmin}_{h \in H} \{E_{in}(h) + \Omega(h)\}$
 - $\Omega(h)$: penalty for complex h
 - Will discuss this when we get to LFD Section 4
- Optimization is a key component in machine learning

Linear Classification

Linear Classification

- Formulation

- Hypothesis set $H = \{h(\vec{x}) = \text{sign}(\vec{w}^T \vec{x})\}$
- Error measure: binary error $e(h(\vec{x}), y) = \mathbb{I}[h(\vec{x}) \neq y]$

- Property

- Simple model (the VC dimension of d-dim perceptron is d+1)
- Good generalization error

- When data is linearly separable

- Run PLA \Rightarrow find g with $E_{in}(g) = 0 \Rightarrow E_{out}(g)$ is close to $E_{in}(g) = 0$

Non-Separable Data

- Generally a hard problem
 - Minimizing E_{in} is a NP-hard problem in general
 - Reason: binary error is discrete and hard to optimize
- Alternative approaches
 - Changing the problem formulation (will discuss in later lectures)
 - Example: Support vector machines in 2nd half of the semester
 - Engineering the features to make data closer to be separable
 - the handwriting digit recognition example
 - Pocket algorithm
 - Run PLA for T rounds
 - Keep track of the best weights \vec{w}^* ($\vec{w}(t)$ that minimizes E_{in})

Linear Regression

Linear Regression

- Formulation
 - Hypothesis set $H = \{h(\vec{x}) = \vec{w}^T \vec{x}\}$
 - Squared error $e(h(\vec{x}), y) = (h(\vec{x}) - y)^2$
- Given dataset $D = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N)\}$
 - $E_{in}(\vec{w}) = \frac{1}{N} \sum_{n=1}^N (\vec{x}_n - y_n)^2$
- Goal: find $\vec{w}_{lin} = \operatorname{argmin}_{\vec{w}} E_{in}(\vec{w})$

Matrix Representation

- $D = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N)\}$

- $X = \begin{bmatrix} \vec{x}_1^T \\ \vdots \\ \vec{x}_N^T \end{bmatrix} = \begin{bmatrix} x_{1,0} & x_{1,1} & \cdots & x_{1,d} \\ x_{2,0} & x_{2,1} & \cdots & x_{2,d} \\ \vdots & \vdots & \ddots & \vdots \\ x_{N,0} & x_{N,1} & \cdots & x_{N,d} \end{bmatrix}$ and $\vec{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$



$x_{n,i}$: the i -th element
of vector \vec{x}_n

Rewriting the In-Sample Error In Matrix Form

$$E_{in}(\vec{w}) = \frac{1}{N} \sum_{n=1}^N (\vec{w}^T \vec{x}_n - y_n)^2$$

$$= \frac{1}{N} \sum_{n=1}^N (\vec{x}_n^T \vec{w} - y_n)^2$$

$$= \frac{1}{N} \|X\vec{w} - \vec{y}\|^2$$

$$= \frac{1}{N} (X\vec{w} - \vec{y})^T (X\vec{w} - \vec{y})$$

$$\|\vec{z}\| = \sqrt{\sum_{i=1}^d z_i^2} = \sqrt{\vec{z}^T \vec{z}}$$



$$E_{in}(\vec{w}) = \frac{1}{N} \left((X\vec{w})^T - \vec{y}^T \right) (X\vec{w} - \vec{y})$$

$$= \frac{1}{N} (\vec{w}^T X^T X \vec{w} - 2\vec{w}^T X^T \vec{y} + \vec{y}^T \vec{y})$$

How to find $\vec{w}_{lin} = \operatorname{argmin}_{\vec{w}} E_{in}(\vec{w})$?

- Answer: Solve for $\nabla_{\vec{w}} E_{in}(\vec{w}) = 0$
- Derivations
 - $E_{in}(\vec{w}) = \frac{1}{N} (\vec{w}^T X^T X \vec{w} - 2 \vec{w}^T X^T \vec{y} + \vec{y}^T \vec{y})$
 - $\nabla_{\vec{w}} E_{in}(\vec{w}) = \frac{1}{N} (2 X^T X \vec{w} - 2 X^T \vec{y})$
 - $\nabla_{\vec{w}} E_{in}(\vec{w}_{lin}) = 0 \implies X^T X \vec{w}_{lin} = 2 X^T \vec{y}$

- $X^T X \vec{w}_{lin} = 2X^T \vec{y}$
 - Two cases:
 - If $X^T X$ is **invertible** (When $N \gg d$, most of the time, it is invertible)
 - $\vec{w}_{lin} = (X^T X)^{-1} X^T \vec{y}$
 - If $X^T X$ is not invertible
 - Requires special handling (See LFD Problem 3.15 for an example)
 - In practice
 - Define X^\dagger as the pseudo-inverse of X
 - when $X^T X$ is invertible, $X^\dagger = (X^T X)^{-1} X^T$
 - When $X^T X$ is not invertible, “handle” it appropriately (usually done in the library for you)
- Linear regression algorithm (a single step algorithm):
 - $\vec{w}_{lin} = X^\dagger \vec{y}$

Discussion

- Special case of **zero-dimensional** space

$$X = \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix} \Rightarrow X^T X = N \Rightarrow (X^T X)^{-1} = 1/N$$

$$\vec{w}_{lin} = (X^T X)^{-1} X^T \vec{y} = \begin{bmatrix} \frac{1}{N} & \dots & \frac{1}{N} \end{bmatrix} \begin{bmatrix} y_1 \\ \vdots \\ y_N \end{bmatrix} = \frac{1}{N} \sum_{n=1}^N y_n$$

Squared error \Rightarrow mean

Discussion

- Linear regression generalizes very well
 - Under mild conditions (See LFD Exercise 3.4 for an example)

$$E_{out}(g) = E_{in}(g) + O\left(\frac{d}{N}\right)$$

- Use regression for classification
 - Note that $\{-1, +1\} \subset \mathbb{R}$
 - Use linear regression to find $\vec{w}_{lin} = (X^T X)^{-1} X^T \vec{y}$ for data with $y \in \{-1, +1\}$
 - Use \vec{w}_{lin} for classification: $g(\vec{x}) = \text{sign}(\vec{w}_{lin}^T \vec{x})$
 - Alternatively, use \vec{w}_{lin} as the initialization for Pocket Algorithm