

CSE417T – Lecture 19

- Please **mute** yourself and **turn off videos** to save bandwidth.
- If you have questions during the lecture
 - Use chatrooms to post your questions
 - I'll review chatrooms in batches
 - You can also un-mute yourself and ask the questions directly
- The slides are posted on the course website
- **RECORD THE LECTURE!**
 - Please remind me if I forget to do so.

Logistics: Homework

- Homework 4 will be due April 13 (Monday)
 - Please start it early
 - It was on average the most time consuming assignment for students in the past
 - Keep track of your own late days
 - Gradescope doesn't allow separate deadlines
 - Your submissions won't be graded if you exceed the late-day limit
- Homework 5 will have a tighter deadline
 - Tentative dates (still subject to change)
 - announce on April 7, due on April 19, 11AM

Recap


Nearest Neighbor

- Predict \vec{x} according to its nearest neighbor
 - Given $D = \{(\vec{x}_1, y_1), (\vec{x}_2, y_2), \dots, (\vec{x}_N, y_N)\}$
 - Let $\vec{x}_{[1]}$ be \vec{x} 's nearest neighbor, i.e., the closest point to \vec{x} in D
 - Let $y_{[i]}(\vec{x})$ or $y_{[i]}$ be the label of $\vec{x}_{[i]}$
- Nearest neighbor hypothesis

$$g(\vec{x}) = y_{[1]}(\vec{x})$$

Nearest Neighbor

$g(\vec{x})$ looks like a Voronoi diagram



- Properties of Nearest Neighbor (NN)
 - No training is needed
 - Good interpretability
 - In-sample error $E_{in} = 0$
 - VC dimension is ∞
- This seems to imply bad learning models from what we talk about so far? Why we care?
- What we really care about is E_{out}
 - VC analysis: $E_{out} \leq E_{in} + \text{Generalization error}$
 - We can infer E_{out} through E_{in} and model complexity
 - NN has nice guarantees outside of VC analysis

Nearest Neighbor is 2-Optimal

- Given mild conditions, for nearest neighbor, when $N \rightarrow \infty$, with high probability,

$$E_{out} \leq 2E_{out}^*$$

- That is, we can not infer E_{out} from E_{in} , but we know it cannot be much worse than the **best anyone can do**.

k -Nearest Neighbor (K-NN)

- Instead of "single" nearest neighbor
 - Making predictions according to k nearest neighbors
- k -nearest neighbor (K-NN)
 - $g(\vec{x}) = \text{sign}(\sum_{i=1}^k y_{[i]}(\vec{x}))$
 - (k is often odd for binary classification)

How to Choose k

- Making the choice of k a function of N , denoted by $k(N)$
 - Theorem:
 - If $k(N) \rightarrow \infty$ as $N \rightarrow \infty$ and $\frac{k(N)}{N} \rightarrow 0$ as $N \rightarrow \infty$
 - Then $E_{in}(g) \rightarrow E_{out}(g)$ and $E_{out}(g) \rightarrow E_{out}(g^*)$
 - Example: $k(N) = \sqrt{N}$ satisfies the condition
- Practical rule of thumb:
 - $k = 3$ is often a good enough choice
 - Using (cross-)validation to choose k

Summary of k -NN

- Pros
 - Simple algorithm
 - Good interpretations
 - Nice theoretical guarantee
 - Easy to adapt to regression (average of nearest neighbors) and multi-class classification (majority voting)
- Cons
 - Computational issue [See LFD 6.2.2 for more discussion]
 - Curse of dimensionality
- We have skipped some more discussion. See LFD Chap 6 if interested.


Lecture Notes Today

The notes are not intended to be comprehensive. They should be accompanied by lectures and/or textbook.
Let me know if you spot errors.

Support Vector Machines (SVM)


What Do We Know about Linear Classification?

- What we discussed so far:
 - PLA: Find a linear separator that separates the data within finite steps, if data is linear separable.
 - Pocket algorithm: empirically keep the best separator during PLA.
 - Surrogate loss: Using logistic regression for linear classification.
- Challenges
 - Binary classification error is hard to optimize,
 - We cannot use “gradient descent” type of algorithm minimize E_{in} .
- Support vector machines (SVM) tries to look at things a bit differently.



Linear Classification


- Which separator would you choose?



Probably the right one.
Why?

Linear Classification


- Which separator would you choose?



More robust to noise (e.g., measurement error of \vec{x})

Linear Classification

- Which separator would you choose?



Margin: shortest distance from the separator to the points in D
(Informal argument)

Higher margin \Rightarrow more “constrained” hypothesis \Rightarrow lower VC dimension

Support Vector Machine

- Goal:
 - Find the **max-margin** linear separator that separates the data
 - Recall the goal of PLA: Find the linear separator that separates the data
- Notations:

Notations we used so far:

- $\vec{x} = (x_0, x_1, \dots, x_d)$
- $\vec{w} = (w_0, w_1, \dots, w_d)$
- Linear separator
$$h(\vec{x}) = \text{sign}(\vec{w}^T \vec{x})$$

Notations we will use in SVM


- $\vec{x} = (x_1, \dots, x_d)$
- $\vec{w} = (w_1, \dots, w_d)$
- Linear separator
$$h(\vec{x}) = \text{sign}(\vec{w}^T \vec{x} + b)$$

Separating the bias/intercept b is important for us to characterize the margin.

We will use (\vec{w}, b) to characterize the hypothesis


Relevant Review of Linear Algebra

- Claim: \vec{w} is the norm vector of the hyperplane $\vec{w}^T \vec{x} + b = 0$



Relevant Review of Linear Algebra


- Claim: \vec{w} is the norm vector of the hyperplane $\vec{w}^T \vec{x} + b = 0$



- Consider any two points \vec{x}' and \vec{x}'' on the hyperplane
 - $\vec{w}^T \vec{x}' + b = 0$
 - $\vec{w}^T \vec{x}'' + b = 0$
- Combining the above
 - $\vec{w}^T (\vec{x}' - \vec{x}'') = 0$
- \vec{w} is orthogonal to the hyperplane
- \vec{w} is the norm vector of the hyperplane


Relevant Review of Linear Algebra

- What is the distance between a point \vec{x}_0 and a hyperplane $\vec{w}^T \vec{x} + b = 0$



Relevant Review of Linear Algebra

- What is the distance between a point \vec{x}_0 and a hyperplane $\vec{w}^T \vec{x} + b = 0$



- Consider an arbitrary point \vec{x}' on the hyperplane
- Distance between the point \vec{x} and the hyperplane



$$\begin{aligned} \text{dist}(\vec{x}_0, \vec{w}, b) &= \left| \frac{\vec{w}^T}{\|\vec{w}\|} (\vec{x}_0 - \vec{x}') \right| \\ &= \left| \frac{1}{\|\vec{w}\|} (\vec{w}^T \vec{x}_0 - \vec{w}^T \vec{x}') \right| \\ &= \left| \frac{1}{\|\vec{w}\|} (\vec{w}^T \vec{x}_0 + b) \right| \end{aligned}$$

Outline of Our Discussion

- Assume data is linearly separable
 - Formulate the **hard-margin SVM**

Given D , find separator (\vec{w}, b) that
maximize $\text{margin}(\vec{w}, b)$
s.t. all points in D is correctly classified

- When data is not linearly separable
 - Tolerate some noise
 - **Soft-margin SVM**
 - Nonlinear transform
 - **Dual formulation** and **kernel tricks**



Hard-Margin SVM

Hard-Margin SVM

- Goal
 - Given $D = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N)\}$ that is **linearly separable**
 - Find separator (\vec{w}, b) that (1) maximizes the margin and (2) separates D
- (\vec{w}, b) separates D (making correct predictions for all points in D)
 - $y_n = \text{sign}(\vec{w}^T \vec{x}_n + b)$ for all n
 - $y_n(\vec{w}^T \vec{x}_n + b) \geq 0$ for all n
- Margin: shortest distance from the separator to points in D

$$\begin{aligned}\text{margin}(\vec{w}, b) &= \min_n \text{dist}(\vec{x}_n, \vec{w}, b) \\ &= \min_n \left| \frac{1}{\|\vec{w}\|} (\vec{w}^T \vec{x}_n + b) \right| \\ &= \min_n \frac{1}{\|\vec{w}\|} y_n (\vec{w}^T \vec{x}_n + b)\end{aligned}$$

$$\text{dist}(\vec{x}_0, \vec{w}, b) = \left| \frac{1}{\|\vec{w}\|} (\vec{w}^T \vec{x}_0 - b) \right|$$

$$y_n \in \{-1, +1\} \text{ and } y_n(\vec{w}^T \vec{x}_n + b) \geq 0$$

Hard-Margin SVM

- Goal
 - Given $D = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N)\}$ that is **linearly separable**
 - Find separator (\vec{w}, b) that (1) maximizes the margin and (2) separates D
- Formulate it as a constrained optimization problem

$$\begin{aligned} &\text{maximize}_{\vec{w}, b} \quad \text{margin}(\vec{w}, b) \\ &\text{subject to} \quad y_n(\vec{w}^T \vec{x}_n + b) \geq 0, \forall n \\ &\quad \text{margin}(\vec{w}, b) = \min_n \frac{1}{\|\vec{w}\|} y_n(\vec{w}^T \vec{x}_n + b) \end{aligned}$$

Hard-Margin SVM

- The constrained optimization problem

$$\begin{aligned} & \text{maximize}_{\vec{w}, b} \quad \text{margin}(\vec{w}, b) \\ & \text{subject to} \quad y_n(\vec{w}^T \vec{x}_n + b) \geq 0, \forall n \\ & \quad \quad \quad \text{margin}(\vec{w}, b) = \min_n \frac{1}{\|\vec{w}\|} y_n(\vec{w}^T \vec{x}_n + b) \end{aligned}$$

- normalizing (\vec{w}, b)
 - Note that $\vec{w}^T \vec{x} + b = 0$ is equivalent to $c\vec{w}^T \vec{x} + cb = 0$ for any c
 - We will normalize (\vec{w}, b) such that $\min_n y_n(\vec{w}^T \vec{x}_n + b) = 1$
 - $y_n(\vec{w}^T \vec{x}_n + b) \geq 1, \forall n$
 - $\text{margin}(\vec{w}, b) = \frac{1}{\|\vec{w}\|}$

Hard-Margin SVM

- The constrained optimization problem

$$\begin{array}{ll} \text{maximize}_{\vec{w}, b} & \frac{1}{\|\vec{w}\|} \\ \text{subject to} & y_n(\vec{w}^T \vec{x}_n + b) \geq 1, \forall n \end{array}$$

- Some final adjustments

$$\begin{array}{ll} \text{minimize}_{\vec{w}, b} & \frac{1}{2} \vec{w}^T \vec{w} \\ \text{subject to} & y_n(\vec{w}^T \vec{x}_n + b) \geq 1, \forall n \end{array}$$

Final Form of Hard-Margin SVM

$$\begin{array}{ll} \text{minimize}_{\vec{w}, b} & \frac{1}{2} \vec{w}^T \vec{w} \\ \text{subject to} & y_n (\vec{w}^T \vec{x}_n + b) \geq 1, \forall n \end{array}$$

- How to solve it?
 - Hard-margin SVM is a Quadratic Program
 - Standard form of Quadratic Program (QP)

$$\begin{array}{ll} \text{minimize}_{\vec{u}} & \frac{1}{2} \vec{u}^T Q \vec{u} + \vec{p}^T \vec{u} \\ \text{subject to} & A \vec{u} \geq \vec{c} \end{array}$$

- There exist efficient QP solvers we can utilize

Short Break and Questions

Connection to Regularization

- Weight decay regularization in linear model

$$\begin{aligned} &\text{minimize } E_{in}(\vec{w}) \\ &\text{subject to } \vec{w}^T \vec{w} \leq C \end{aligned}$$


- Another way to look at SVM

$$\begin{aligned} &\text{minimize } \vec{w}^T \vec{w} \\ &\text{subject to } E_{in}(\vec{w}) = 0 \end{aligned}$$

Support Vectors

- We call the points closest to the separator (candidate) **support vectors**
 - Since they **support** the separator
- What are the math properties of support vectors?
 - They are the points that the equality holds in the constraints
 - If \vec{x}_n is a support vector, $y_n(\vec{w}^T \vec{x}_n + b) = 1$


$$\begin{array}{ll} \text{minimize}_{\vec{w}, b} & \frac{1}{2} \vec{w}^T \vec{w} \\ \text{subject to} & y_n(\vec{w}^T \vec{x}_n + b) \geq 1, \forall n \end{array}$$



- Removing the non-support vectors will not impact the linear separator

Leave-One-Out Cross Validation (LOOCV)


- Two things we know so far
 - Removing non-support vectors will not impact the separator
 - LOOCV error (when not used for model selection) is an unbiased estimate of $E_{out}(N - 1)$ (E_{out} when trained on $N - 1$ points)
- What's the LOOCV error for SVM?
 - $E_{LOOCV} \leq \frac{\# \text{ support vectors}}{N}$
- Note that we know # support vectors after training
 - Count # points that satisfy $y_n(\vec{w}^T \vec{x}_n + b) = 1$
- Another method to estimate/bound E_{out} (counting # support vectors)




What if Data is Not Linearly Separable

Non-Separable Data

- Two scenarios




- Tolerate some noise
 - **Soft-Margin SVM**



- Nonlinear transform
 - **Dual formulation and kernel tricks**

Soft-Margin SVM


- Intuition: We want to tolerate small noises when maintaining large margin
- For each point (\vec{x}_n, y_n) , we allow a deviation $\xi_n \geq 0$
 - Instead of requiring $y_n(\vec{w}^T \vec{x}_n + b) \geq 1$
 - The constraint becomes
$$y_n(\vec{w}^T \vec{x}_n + b) \geq 1 - \xi_n$$
- We add a penalty for each deviation
 - Total penalty $C \sum_{n=1}^N \xi_n$



Soft-Margin SVM

- The constraint becomes: $y_n(\vec{w}^T \vec{x}_n + b) \geq 1 - \xi_n$
- We add a penalty for each deviation: Total penalty $C \sum_{n=1}^N \xi_n$

$$\begin{array}{ll} \text{minimize}_{\vec{w}, b, \vec{\xi}} & \frac{1}{2} \vec{w}^T \vec{w} + C \sum_{n=1}^N \xi_n \\ \text{subject to} & y_n(\vec{w}^T \vec{x}_n + b) \geq 1 - \xi_n, \forall n \\ & \xi_n \geq 0, \forall n \end{array}$$




Remarks:


- C is a hyper-parameter we can choose, e.g., using validation
- Soft-margin SVM is still a Quadratic Program, with efficient solvers

Impacts of C in Soft-Margin SVM

- Large C => less tolerate to noise, having smaller margin




$C = 1$



$C = 500$

What if Tolerating Small Noises Is Not Enough



Nonlinear transform

We can apply standard nonlinear transformation procedure we talked about before

In SVM, we can combine the ideas of **dual formulation** and **kernel tricks** for the transformation

This is one of the key ingredients that makes SVM powerful