

CSE417T – Lecture 23

- Please **mute** yourself and **turn off videos** to save bandwidth.
- If you have questions during the lecture
 - Use chatrooms to post your questions
 - I'll review chatrooms in batches
 - You can also un-mute yourself and ask the questions directly
- The slides are posted on the course website
- **RECORD THE LECTURE!**
 - Please remind me if I forget to do so.

Logistics: Homework and Exam 2

- Homework 4 was due Monday.
 - Latest submission by today if you use 3 late days
- Homework 5 will be due on April 19 (Sunday), **11:30AM**
 - At most two late days can be used in this homework
 - I'll post the submission link on Friday to avoid conflicts with HW4 submissions
- Test exam is on Canvas till **Friday 1pm**
 - Strongly encouraged to do it make sure the setup works in your environment
 - No additional extension will be given for exam 2 due to unfamiliarity of the tool
- Exam 2 will be online on Canvas on April 23 (Thursday)
 - See Slides on April 7 for more details

Logistics: Exam 2

- Exam 2 will be on April 23 using **Canvas Quiz + Lockdown Browser**.
- Exam duration: **80 minutes**
 - 5 more minutes than Exam 1 as the buffer for online exam
- Start time
 - **11:30am CDT** (lecture start time).
 - If you cannot make it, please let me know by **tomorrow, April 17**
 - Only students who get approved can take the exam at a different time.
 - Unless there is a strong reason, everyone should take the exam on April 23.

Logistics: Exam 2

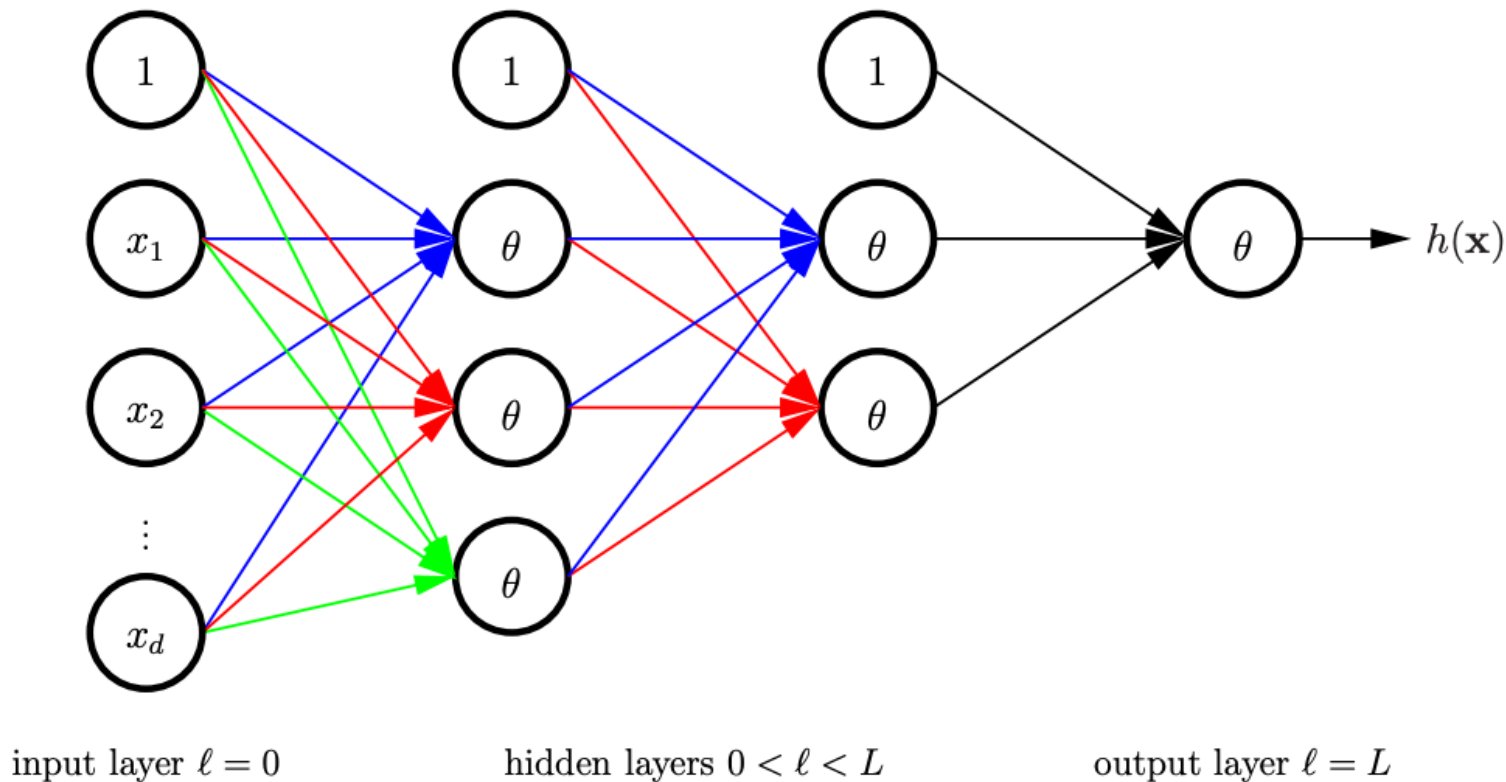
- Format
 - A mix of long questions and multiple choice questions
 - Likely dominated by multiple choices due to the constraint of math typing
 - I will try to minimize the need to write math in the long questions.
- Open book
 - You can reference any materials in hard copies. Searching information online is not allowed. Talking with other people is not allowed.
 - LockDown Browser will prohibit you from accessing other information in your computer
- Randomized questions
 - The questions will be randomly drawn from a “question bank”, so everyone might be getting different questions. I’ll take that into account in final grades.

Logistics: Exam 2

- Topic
 - The focus will be on the 2nd half of the semester
 - Knowledge is cumulative, and concepts in Exam 1 might also be included
 - Everything in lectures and in readings on the course website is included (except for the parts marked as [safe to skip](#)).
- Remaining lectures
 - ~~Apr 14 (Tue): Learning in Neural Networks~~
 - Apr 16 (Thu): Discussion on Neural Networks
 - Apr 21 (Tue): Brief review and office hour
 - Apr 23 (Thu): Exam 2

Recap

Neural Networks





θ : **activation function**
(Specify the “activation” of the neuron)



We mostly focus on **feed-forward** network structure

Notations of Neural Networks (NN)

- Notations:
 - $\ell = 0$ to L : layer
 - $d^{(\ell)}$: dimension of layer ℓ
 - $\vec{x}^{(\ell)}$: the nodes in layer ℓ
 - $w_{i,j}^{(\ell)}$: weights; characterize hypothesis in NN
 - $s_j^{(\ell)} = \sum_{i=0}^{d^{(\ell-1)}} w_{i,j}^{(\ell)} x_i^{(\ell-1)}$: linear signals
 - θ : activation function
 - $x_j^{(\ell)} = \theta(s_j^{(\ell)})$



Forward Propagation (evaluate $h(\vec{x})$)

- A Neural network hypothesis h is characterized by $\{w_{i,j}^{(\ell)}\}$
- How to evaluate $h(\vec{x})$?

$$\mathbf{x} = \mathbf{x}^{(0)} \xrightarrow{w^{(1)}} \mathbf{s}^{(1)} \xrightarrow{\theta} \mathbf{x}^{(1)} \xrightarrow{w^{(2)}} \mathbf{s}^{(2)} \xrightarrow{\theta} \mathbf{x}^{(2)} \dots \xrightarrow{w^{(L)}} \mathbf{s}^{(L)} \xrightarrow{\theta} \mathbf{x}^{(L)} = h(\mathbf{x}).$$

Forward propagation to compute $h(\mathbf{x})$:

1: $\mathbf{x}^{(0)} \leftarrow \mathbf{x}$	[Initialization]
2: for $\ell = 1$ to L do	[Forward Propagation]
3: $\mathbf{s}^{(\ell)} \leftarrow (W^{(\ell)})^T \mathbf{x}^{(\ell-1)}$	
4: $\mathbf{x}^{(\ell)} \leftarrow \begin{bmatrix} 1 \\ \theta(\mathbf{s}^{(\ell)}) \end{bmatrix}$	
5: end for	
6: $h(\mathbf{x}) = \mathbf{x}^{(L)}$	[Output]

Given weights $w_{i,j}^{(\ell)}$ and $\vec{x}^{(0)} = \vec{x}$, we can calculate all $\vec{x}^{(\ell)}$ and $\vec{s}^{(\ell)}$ through forward propagation.

How to Learn NN From Data?

- Given D , how to learn the weights $W = \{w_{i,j}^{(\ell)}\}$?
- Intuition: Minimize $E_{in}(W) = \frac{1}{N} \sum_{n=1}^N e_n(W)$
- How?
 - Gradient descent: $W(t+1) \leftarrow W(t) - \eta \nabla_W E_{in}(W)$
 - Stochastic gradient descent $W(t+1) \leftarrow W(t) - \eta \nabla_W e_n(W)$
- Key step: we need to be able to evaluate the gradient...
 - Not trivial to do given the network structure
 - **Backpropagation** is an algorithmic procedure to calculate the gradient

Compute the Gradient $\nabla_W e_n(W)$

- Applying chain rule

$$\frac{\partial e_n(W)}{\partial w_{i,j}^{(\ell)}} = \frac{\partial e_n(W)}{\partial s_j^{(\ell)}} \frac{\partial s_j^{(\ell)}}{\partial w_{i,j}^{(\ell)}} = \delta_j^{(\ell)} x_i^{(\ell-1)}$$

- Calculating $\delta_j^{(\ell)}$

- Backward recursive formulation



$$\delta_j^{(\ell)} = \sum_{k=1}^{d^{(\ell+1)}} \frac{\partial e_n(W)}{\partial s_k^{(\ell+1)}} \frac{\partial s_k^{(\ell+1)}}{\partial x_j^{(\ell)}} \frac{\partial x_j^{(\ell)}}{\partial s_j^{(\ell)}} = \sum_{k=1}^{d^{(\ell+1)}} \delta_k^{(\ell+1)} w_{j,k}^{(\ell+1)} \theta' \left(s_j^{(\ell)} \right)$$

- Boundary conditions

- The output layer (assume regression)

$$\delta_1^{(L)} = 2 \left(s_1^{(L)} - y_n \right) \text{ (generalizable to other differentiable error)}$$

- Backward propagation



Backpropagation Algorithm

- Recall that $\frac{\partial e_n(W)}{\partial w_{i,j}^{(\ell)}} = \delta_j^{(\ell)} x_i^{(\ell-1)}$
- Backpropagation Algorithm
 - Initialize $w_{i,j}^{(\ell)}$ randomly
 - For $t = 1$ to T
 - Randomly pick a point from D (for stochastic gradient descent)
 - Forward propagation: Calculate all $x_i^{(\ell)}$ and $s_i^{(\ell)}$
 - Backward propagation: Calculate all $\delta_j^{(\ell)}$
 - Update the weights $w_{i,j}^{(\ell)} \leftarrow w_{i,j}^{(\ell)} - \eta \delta_j^{(\ell)} x_i^{(\ell-1)}$
 - Return the weights

Discussion

- Backpropagation is gradient descent with efficient gradient computation
- Note that the E_{in} is **not convex** in weights
- Gradient descent doesn't guarantee to converge to global optimal
- Common approaches:
 - Run it many times
 - Each with a different initialization (the choice of initialization matters)

Lecture Notes Today

The notes are not intended to be comprehensive. They should be accompanied by lectures and/or textbook.
Let me know if you spot errors.

Regularization in Neural Networks

Weight-Based Regularization

- Weight decay

$$E_{aug}(W) = E_{in}(W) + \frac{\lambda}{N} \sum_{i,j,\ell} \left(w_{i,j}^{(\ell)} \right)^2$$

- Backpropagation applies (changing the error measure)



- Weight elimination

$$E_{aug}(W) = E_{in}(W) + \frac{\lambda}{N} \sum_{i,j,\ell} \frac{\left(w_{i,j}^{(\ell)} \right)^2}{1 + \left(w_{i,j}^{(\ell)} \right)^2}$$

- When $w_{i,j}^{(\ell)}$ is small, approximates weight decay
- When $w_{i,j}^{(\ell)}$ is large, approximates adding a constant (no impacts to gradient)
- “Decaying” more on smaller weights (i.e., eliminating small weights)



Early Stopping

- Consider gradient descent (GD)
 - H_1 : the set of hypothesis GD can reach at $t = 1$
 - H_T : the set of hypothesis GD can reach at $t = T$
 - $H_1 \subseteq H_2 \subseteq H_3 \subseteq \dots$




Early Stopping

- Stopping gradient descent early is a regularization method
 - “Constrained the hypothesis set”
- How to find the optimal stopping point t^* ?
 - Using validation is a common approach



Dropout



- Neural networks is very expressive (low bias, potentially high variance)
- Dropout
 - Randomly **drop** p portion of the weights during training



- Learn many models with dropout
- **Average** them during prediction (reduce weights by a ratio of p)


A Nontraditional Method to Avoid Overfitting

- What's the cause of overfitting?



- Fitting the **noise** instead of the target
- Regularization: Constrain H so it's not that powerful to fit noise
- How about **adding noises** to data?

Adding Noises as Regularization



Short Break and Q&A

Deep Learning

Single Hidden-Layer Neural Network

- How do we write a hypothesis in single-hidden layer mathematically?

- $$h(\vec{x}) = \theta \left(w_{0,1}^{(2)} + \sum_{j=1}^{d^{(1)}} w_{j,1}^{(2)} x_{j,1}^{(1)} \right)$$
$$= \theta \left(w_{0,1}^{(2)} + \sum_{j=1}^{d^{(1)}} w_{j,1}^{(2)} \theta \left(\sum_{i=0}^{d^{(0)}} w_{i,j}^{(1)} x_i \right) \right)$$

- How do we write a Kernel SVM hypothesis


- $$g(\vec{x}) = \theta \left(b^* + \sum_{\alpha_n^* > 0} \alpha_n^* y_n K(\vec{x}_n, \vec{x}) \right)$$

- How do we write a linear model with nonlinear transform

- $$h(\vec{x}) = \theta(w_0 + \sum w_i \phi_i(\vec{x}))$$

- Interpretation:

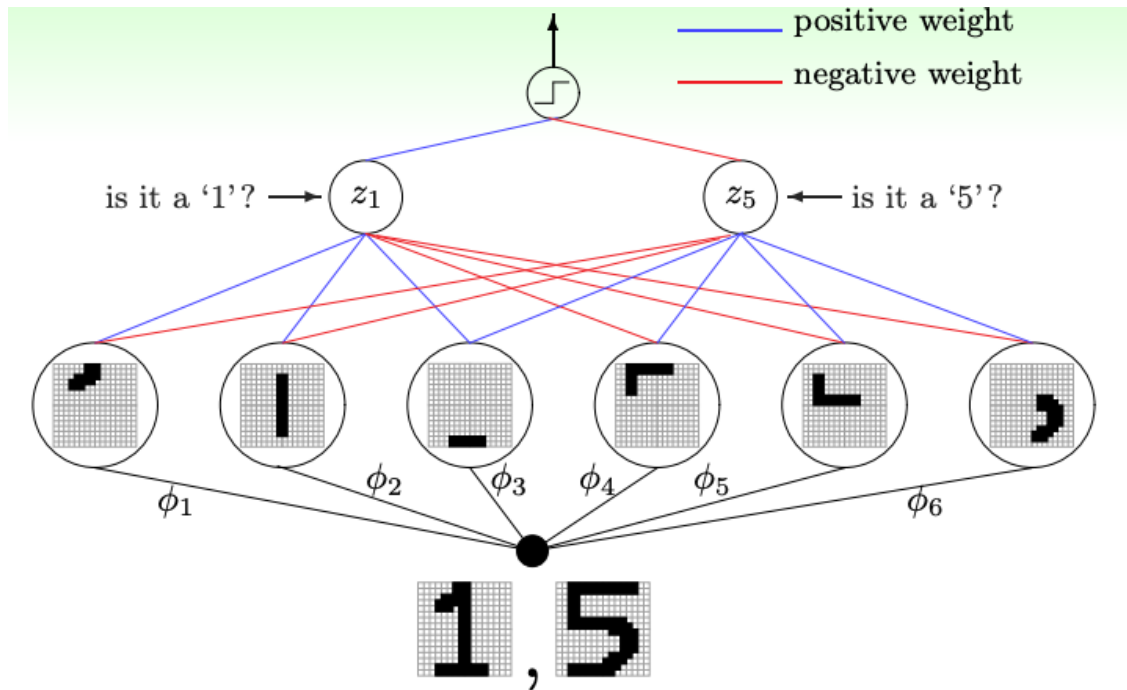
- The hidden layer is like **feature transform**
 - Shallow learning vs. deep learning



[Safe to Skip Starting this Slide]

Deep Neural Network

- “Shallow” neural network is powerful (universal approximation theorem holds with a single hidden layer). Why “deep” neural networks?




Each layer captures **features** of the previous layers.

We can use “raw data” (e.g., pixels of an image) as input. The hidden layer are extracting the **features**.

Design different **network architectures** to incorporate domain knowledge.

Convolutional Neural Networks

- Captures the localized properties of features
 - Particularly suitable for computer vision (images)
 - Go (AlphaGo) is another famous application of CNN




Some Techniques in Improving Deep Learning

- Regularization to mitigate overfitting
 - Weight-based, early stopping, dropout, etc
- Incorporating domain knowledges
 - Network architectures (e.g., Convolutional Neural Nets)
- Improving computation with huge amount of data
 - Hardware architecture to improve parallel computation
- Improving gradient-based optimization
 - Choosing better **initialization** points

Initialization

Minimizing Error is Nonconvex in Deep Learning



- We mostly adopt gradient-descent-style algorithms for optimization.
- No guarantee to converge to global optimal.
- Need to run it many times.
- Initialization matters!

Vanishing Gradient Problem

- Backpropagation

- $\frac{\partial e_n(W)}{\partial w_{i,j}^{(\ell)}} = \delta_j^{(\ell)} x_i^{(\ell-1)}$


- $\delta_j^{(\ell)} = \theta' \left(s_j^{(\ell)} \right) \sum_{k=1}^{d^{(\ell+1)}} \delta_k^{(\ell+1)} w_{j,k}^{(\ell+1)}$

- If we use activation function $\theta(s) = \tanh(s)$

- $\theta'(s) = 1 - \theta(s)^2 < 1$


- In deep learning with a lot of layers,

- the gradient will vanish
 - hard to update the early layers




Vanishing Gradient Problem

- $\delta_j^{(\ell)} = \theta' \left(s_j^{(\ell)} \right) \sum_{k=1}^{d^{(\ell+1)}} \delta_k^{(\ell+1)} w_{j,k}^{(\ell+1)}$
- There is also a corresponding “exploding gradient problem”
- What can we do
 - Choose more suitable activation functions
 - One common choice is Rectified Linear Unit (ReLU) and its variant
 - $\theta(s) = \max(0, s)$
 - Choose better **initialization**




Initialization

- Hard to initialize the entire network well.
- Intuition: Initialize the weights **layer by layer** such that each layer **preserves** the properties of the previous layer.



Autoencoder

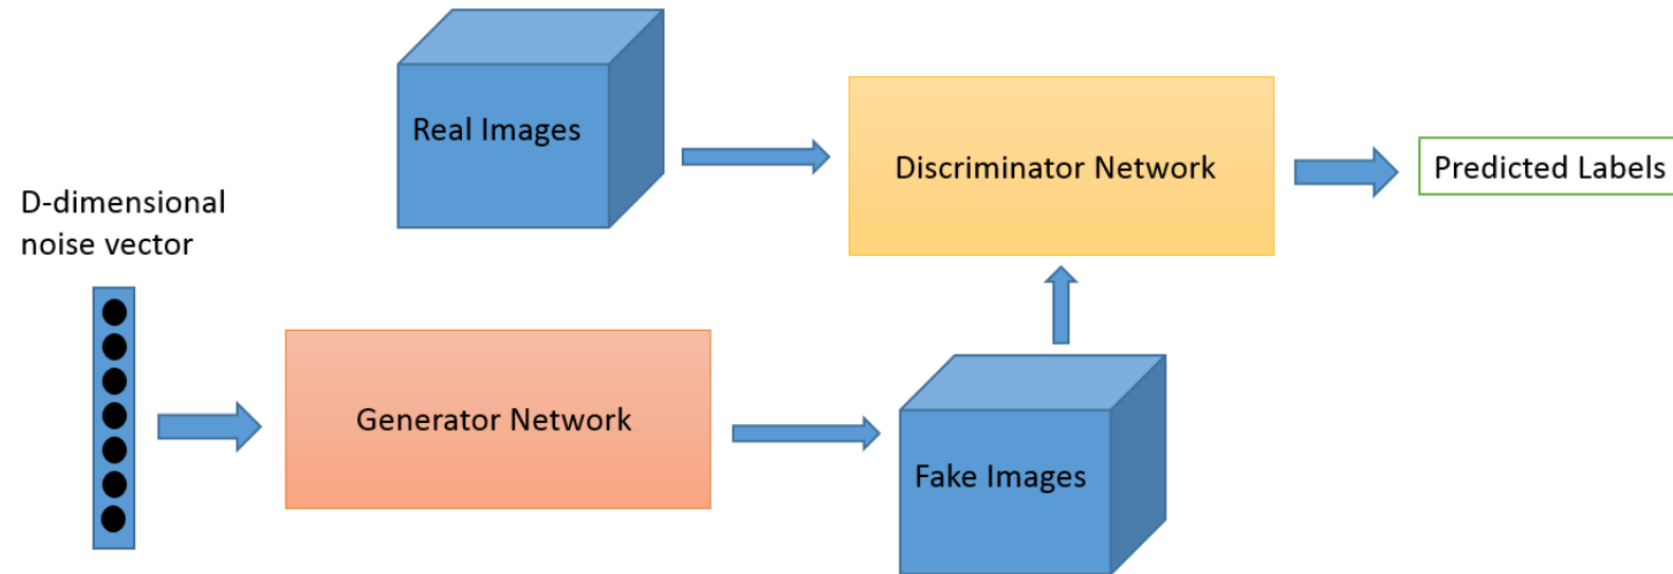


Unsupervised learning!

Generative Adversarial Nets (GAN)

A Competition: Generator vs Discriminator

- Discriminator wants to correctly classify the images (true images or not)
- Generator wants to generate images that discriminator can't classify



Credit: O'Reilly



<https://thisPersonDoesNotExist.com/>